

HW 09. Linear Regression

We continue to work with the same [SOCCR](#) data source as the last time.

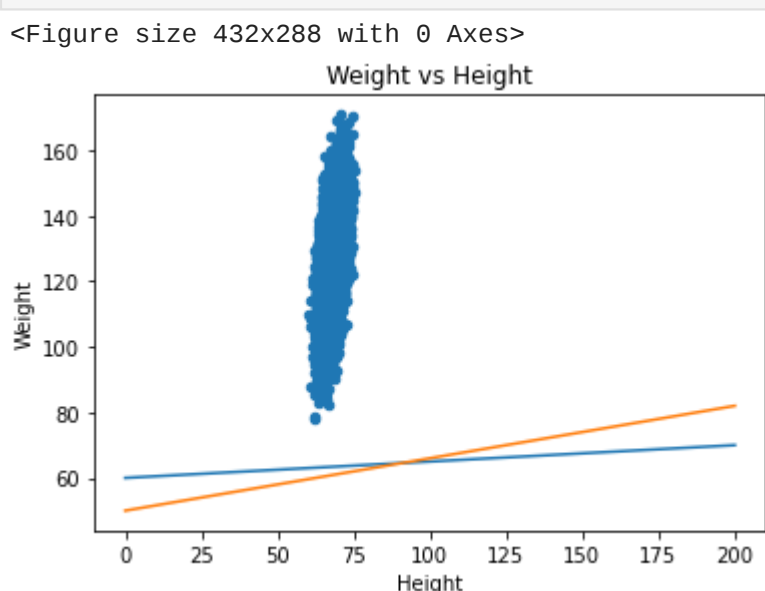
```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.optimize import minimize_scalar
from scipy.optimize import minimize
%matplotlib inline
```

Import the data to a dataframe:

```
In [2]: data = pd.read_csv('teenagers.csv', index_col='Index')
```

Task 1: Use the method `plot()` from the module `matplotlib.pyplot` to plot the dependence of the weight on the height from the dataset. Add the plots of the two linear regressions to the graph: one with the coefficients $(w_0, w_1) = (60, 0.05)$ and the other with $(w_0, w_1) = (50, 0.16)$. Do the regressions approximate the data well? Don't forget to label the axis and add a title to the figure.

```
In [48]: fig = plt.figure()
xx = np.linspace(0, 200, 200)
data.plot(x = 'Height', y = 'Weight', kind = 'scatter', title = 'Weight vs Height')
plt.xlabel('Height')
plt.ylabel('Weight')
w0_1, w1_1 = 60, .05
w0_2, w1_2 = 50, .16
plt.plot(xx, w1_1 * xx + w0_1)
plt.plot(xx, w1_2 * xx + w0_2)
plt.show()
#no these regressions do not approximate the data well at all
```




Task 2: Write a python function `error(X, Y, w0, w1)`, which calculates the Loss Function $Q(w_0, w_1) = \sum_{i=1}^n (y_i - (w_0 + w_1 * x_i))^2$

```
In [30]: def error(x, y, w0, w1):
    retval = 0
    errarray = (y - (w0 + w1 * x)) ** 2
    retval = sum(errarray)
    return retval
```

Task 3: Plot the loss function vs w_1 with a fixed value of $w_0 = 50$ for your dataset. Where approximately does the function reaches its minimum?

```
In [33]: fig = plt.figure()
ax = plt.axes()
errors_len = 200
ww1 = np.linspace(-500, 500, errors_len)
ww0 = np.zeros(errors_len) + 50
errors_arr = np.zeros(errors_len)
for i in range(0, errors_len):
    errors_arr[i] = error(data['Height'], data['Weight'], ww0[i], ww1[i])
ax.plot(ww1, errors_arr)
plt.show()
```



Task 4: Use the method `minimize_scalar` from the module `scipy.optimize` to find the minimum of your Loss function for a fixed value of $w_0 = 50$. Make the method to look for the minimum in the range $-5 \leq w_1 \leq 5$.

```
In [47]: def err50(w1):
    return error(data['Height'], data['Weight'], 50, w1)


w1_opt = minimize_scalar(err50, bounds = (-100, 300))
w1_opt.x
```

Out[47]: 1.1351597092091665

Task 5: Use the method `plot()` from the module `matplotlib.pyplot` to plot the dependence of the weight on the height from the dataset. Add to the graph a plot of the linear regression with $w_0 = 50$, and $w_0 =$ the value you found in Task 4. Does this linear regression model the data better than the one from the Task 1?

```
In [51]: fig = plt.figure()
xx = np.linspace(0, 200, 200)
plt.plot(data['Height'], data['Weight'], 'o')
plt.xlabel('Height')
plt.ylabel('Weight')
w0_1, w1_1 = 50, w1_opt.x
plt.plot(xx, w1_1 * xx + w0_1)
plt.show()

#this model is slightly better
```



Task 6: Make a 3D plot of the Loss function vs both w_0 and w_1 . Label the axis and add a title to the plot. According to the plot, where should we look for the minimum of the loss function?

```
In [46]: from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.gca(projection = '3d')

X = np.linspace(-10000, 10000, 40)
Y = np.linspace(-200, 200, 40)


Z = np.zeros((len(X), len(Y)))

def err2d(w0, w1):
    return error(data['Height'], data['Weight'], w0, w1)

for i in range(0, len(X)):
    for j in range(0, len(Y)):
        Z[i, j] = err2d(X[i], Y[j])

X, Y = np.meshgrid(X, Y)

surf = ax.plot_surface(X, Y, Z)
ax.set_xlabel('intercept')
ax.set_ylabel('slope')
ax.set_zlabel('Z')
plt.show()
```



Task 7: Use the method `minimize` from the module `scipy.optimize` to find the minimum of the loss function. Make the method to look for the minimum inside the range $-100 \leq w_0 \leq 100$, $-5 \leq w_1 \leq 5$. Use the argument `method='L-BFGS-B'`


```
In [53]: def error2D(w):
    return error(data['Height'], data['Weight'], w[0], w[1])
res = minimize(error2D, (0,0), method = 'L-BFGS-B')
res.x
```

Out[53]: array([-82.60496837, 3.08391066])

Task 8: Use the method `plot()` from the module `matplotlib.pyplot` to plot the dependence of the weight on the height from the dataset. Add to the graph a plot of the linear regression with $w_0 = w_0$ = the value you found in Task 7, and $w_0 =$ the value you found in Task 4.

```
In [54]: fig1 = plt.figure()
xx = np.linspace(0, 250, 200)
plt.plot(data['Height'], data['Weight'], 'o')
plt.xlabel('Height')
plt.ylabel('Weight')
w0_1, w1_1 = res.x[0], res.x[1]
plt.plot(xx, w1_1 * xx + w0_1)
plt.show()

#by far the best model
```



In []: